

Albert Query - Agentic Movie Intelligence System

An intelligent conversational agent for querying movie and TV series data using Retrieval Augmented Generation (RAG), multi-tool orchestration, and semantic search.

Python 3.8+

Streamlit 1.51.0

LangChain 1.0.3

LangGraph 1.0.2

ChromaDB 1.3.4

Table of Contents

- [Overview](#)
- [Features](#)
- [Architecture](#)
- [Installation](#)
- [Usage](#)
- [Project Structure](#)
- [Components](#)
- [Workflow](#)
- [Technologies](#)
- [Future Improvements](#)
- [Contributors](#)
- [License](#)

Overview

Albert Query is an agentic AI system we developed as part of our M1 project at **Albert School** in collaboration with **Mines Paris - PSL**. The system intelligently answers questions about movies and TV series by orchestrating multiple data sources and tools through a LangGraph-based workflow.

What Makes It Special?

Unlike traditional chatbots, Albert Query:

- **Plans before acting** - Analyzes each question to determine which tools are needed
- **Multi-source intelligence** - Combines SQL databases, vector search, external APIs, and web search
- **Semantic understanding** - Uses OpenAI embeddings to find movies by plot similarity
- **Source attribution** - Always shows where information comes from
- **Context-aware** - Maintains conversation history for follow-up questions

Use Cases

-  **Semantic Search:** "Find me movies about space exploration with AI themes"

-  **Data Analysis:** "How many comedies were released on Netflix after 2020?"
 -  **Movie Discovery:** "Show me films similar to Inception"
 -  **Trend Analysis:** "What are the top-rated action movies from the 2010s?"
 -  **Latest Info:** "What's trending in movies this week?"
-

◆ Features

Core Capabilities

Intelligent Query Planning

- LLM-based planner analyzes questions and conversation history
- Automatically selects optimal tools (SQL, Semantic Search, OMDB, Web)
- Avoids unnecessary API calls for efficiency

Multi-Database SQL Queries

- Comprehensive catalog of 8,000+ movies/shows from Netflix, Disney+, Amazon Prime
- Structured queries with filters (year, genre, rating, type)
- Automatic database schema understanding

Semantic Vector Search

- 114MB of OpenAI embeddings (text-embedding-3-small)
- Find movies by plot descriptions, themes, or similarity
- Natural language queries in English or French

OMDB API Integration

- Enriched movie metadata (actors, awards, ratings, posters)
- IMDb links for detailed information
- Full plot summaries

Web Search

- DuckDuckGo integration for trending topics
- Latest movie news and releases
- Current events in cinema

Source Attribution

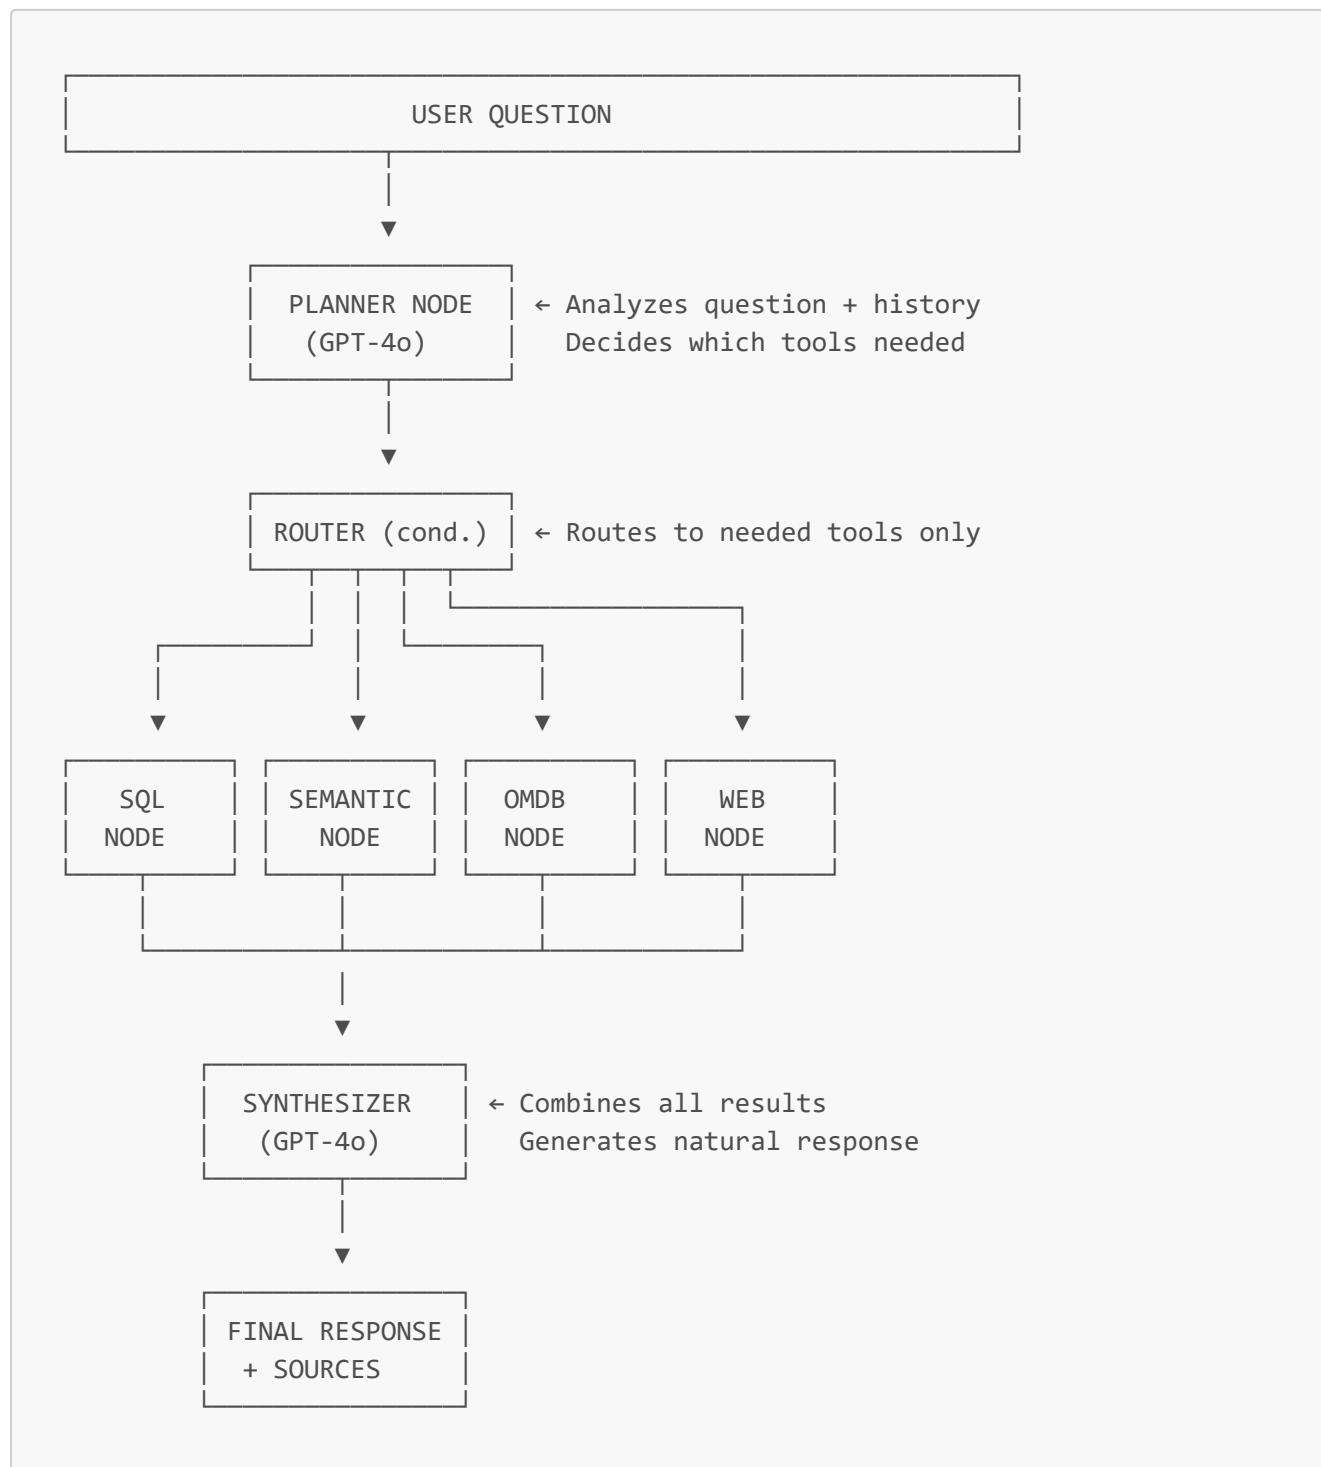
- Tracks all information sources
- Visual display of databases, APIs, and search results used
- Clickable links to IMDb and search results

Conversational Memory

- LangGraph checkpointing with MemorySaver
- Full conversation history maintained
- Context resolution ("that movie" → actual movie name)

E Architecture

Our system follows an **agentic architecture** using LangGraph to create a stateful, multi-tool workflow:



F Project Structure

```
Agentic_Systems_with_RAG_Lamy-Waerniers/
  code/
    albert_query_app.py # Main source code
                           # Main Streamlit application (1009 lines)
```

```

|   └── embedding_manager.py           # Vector embedding management (418
lines)
|   └── architecture.md              # System architecture
documentation
    ├── embeding.ipynb               # Embedding pipeline notebook
    ├── SQLdb_creator.ipynb          # Database creation from CSVs
    ├── testing.ipynb                # API and integration tests
    └── test_semantic_search.ipynb    # Semantic search validation

└── data/
    ├── csv_db/
        ├── amazon_prime_titles.csv   # Source CSV files
        ├── netflix_titles.csv       # Amazon Prime catalog (3.9MB)
        └── disney_plus_titles.csv    # Netflix catalog (3.4MB)
        # Disney+ catalog (385KB)

    ├── databases/
        └── movie.db                 # Consolidated SQLite DB (32.9MB)

    ├── vector_database/
        ├── chroma.sqlite3            # ChromaDB persistent storage
        └── 19c0759d-.../             # Vector DB metadata (42.7MB)

    └── memory/
        ├── conversations/          # Embedding data (114MB)
        └── user_profiles/           # Conversation storage

└── doc/
    ├── graph_schema.png             # Documentation
    ├── omdb_api_doc.json           # LangGraph workflow diagram
    └── OMDB_API_doc.txt            # OMDB API reference

└── .env                           # Environment configuration (git-
ignored)
└── .gitignore                     # Git ignore rules
└── requirements.txt               # Python dependencies (223
packages)
└── README.md                      # This file

```

❖ Components

1. Main Application ([albert_query_app.py](#))

The core orchestrator built with **Streamlit** and **LangGraph**.

Key Components:

Planner Node

```

def planner_node(state: AgentState) -> dict:
    """Analyzes question and plans tool usage"""
    # Uses GPT-4o-mini to decide which tools are needed

```

```
# Returns: needs_sql, needs_semantic, needs_omdb, needs_web flags  
# Also prepares specific queries for each tool
```

Decision Logic:

- SQL: For structured data (titles, years, ratings, genres)
- Semantic: For plot-based searches and "movies like X"
- OMDB: For detailed metadata (actors, awards, posters)
- Web: For trending topics and recent news

SQL Node

```
def sql_node(state: AgentState) -> dict:  
    """Generates and executes SQL queries"""  
    # Analyzes database catalog  
    # Generates SQL with proper table names, filters  
    # Executes query and returns results
```

Features:

- Auto-detects correct table names
- Handles genre filtering (comma-separated values)
- Year ranges, type filtering (Movie vs TV Show)
- Result limiting and ordering

Semantic Search Node

```
def semantic_search_node(state: AgentState) -> dict:  
    """Performs vector similarity search"""  
    # Translates query to English if needed  
    # Optimizes query for embeddings  
    # Searches ChromaDB with OpenAI embeddings  
    # Returns top-k similar movies
```

Important Notes:

- Movie descriptions are in **English** - queries are auto-translated
- Similarity scores may appear low (40-60%) but rankings are accurate
- Uses cosine similarity on text-embedding-3-small vectors

OMDB Node

```
def omdb_node(state: AgentState) -> dict:  
    """Fetches enriched movie data from OMDB API"""
```

```
# Extracts movie title from context
# Queries OMDB API
# Returns: plot, actors, awards, IMDb rating, poster URL
```

Web Search Node

```
def web_node(state: AgentState) -> dict:
    """Performs web search via DuckDuckGo"""
    # Executes search query
    # Returns recent results with URLs
```

Synthesizer Node

```
def synthesizer_node(state: AgentState) -> dict:
    """Combines all results into natural response"""
    # Takes SQL, semantic, OMDB, web results
    # Generates coherent natural language answer
    # Includes source attribution
```

2. Embedding Manager (`embedding_manager.py`)

Handles creation and management of vector embeddings for semantic search.

Key Functions:

```
def get_or_create_collection(chroma_path,
collection_name="movie_descriptions"):
    """Gets or creates ChromaDB collection with OpenAI embeddings"""

def embed_movies_if_not_exists(collection, movies, batch_size=100):
    """Embeds movies only if not already in collection"""
    # Checks existing IDs to avoid duplicates
    # Batches for efficiency
    # Tracks progress with tqdm

def query_movies(collection, query_text, n_results=5, where_filter=None):
    """Queries movies by semantic similarity"""
    # Embeds query with same model
    # Searches vector DB
    # Returns ranked results with similarity scores
```

Embedding Structure:

```
{
    "id": "net0042", # Unique ID (prefix + index)
    "document": "A sci-fi thriller...", # Movie description (embedded)
    "metadata": {
        "title": "The Matrix",
        "database": "movie.db",
        "table": "netflix_titles"
    }
}
```

3. Database Schema

SQL Tables:

Each table (netflix_titles, amazon_prime_titles, disney_plus_titles) contains:

Column	Type	Description
show_id	TEXT	Unique identifier (e.g., "s1")
type	TEXT	"Movie" or "TV Show"
title	TEXT	Movie/show title
director	TEXT	Director name(s)
cast	TEXT	Comma-separated actors
country	TEXT	Production country
date_added	TEXT	Date added to platform
release_year	INTEGER	Year of release
rating	TEXT	Age rating (PG, R, etc.)
duration	TEXT	Runtime or seasons
listed_in	TEXT	Comma-separated genres
description	TEXT	Plot summary

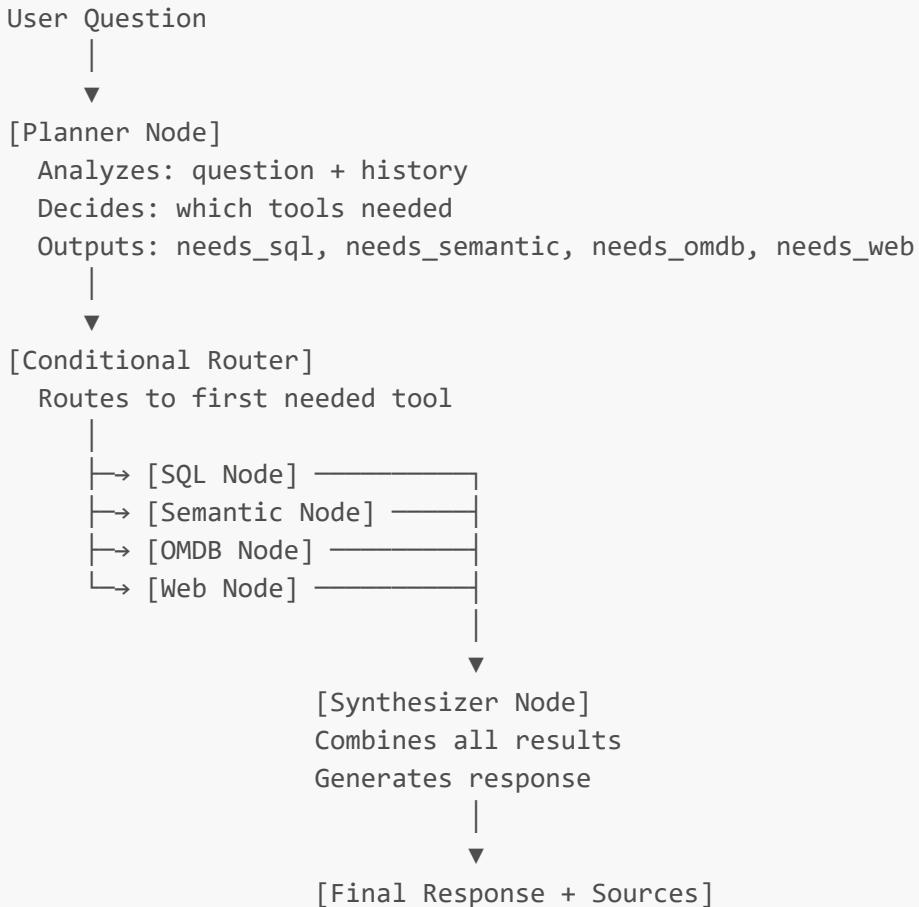
Total Records: ~8,800 movies/shows

4. Vector Database

ChromaDB Collection:

- Name: `movie_descriptions`
- Embedding Model: OpenAI text-embedding-3-small (1536 dimensions)
- Total Vectors: ~8,800 movie descriptions
- Storage: Persistent on disk (~156MB total)
- Metadata: title, database, table for filtering

⌚ Workflow



🚧 Future Improvements

We've identified several areas for optimization and enhancement. Here's our roadmap:

1. 📁 Catalog Caching System

Problem:

- Database catalog is rebuilt on every app startup
- Slow initialization (~2-5 seconds)
- Redundant SQL queries for schema introspection

Solution:

```
# Implement caching with invalidation detection
def get_or_build_catalog(db_path: str, cache_path: str) -> dict:
    """
    Cache database catalog as JSON
    - Compare file modification times to detect changes
    - Load from cache if DB unchanged
    - Rebuild only when necessary
    """
    pass
```

Expected Impact:

- ⚡ 10-50x faster startup time
- 💾 Reduced SQL queries
- ⌚ Auto-invalidation on schema changes

Implementation:

- Save catalog to `data/databases/catalog_cache.json`
- Include DB file mtime and size for change detection
- Add force-rebuild option for manual invalidation

2. 🧠 Persistent Long-Term Memory

Current State:

- Memory stored in LangGraph's MemorySaver (in-memory only)
- Lost on application restart
- No cross-session learning

Proposed Architecture:

```
# SQLite-based conversation storage
conversations/
    └── user_123/
        ├── session_20250116_001.json      # Conversation history
        ├── session_20250116_002.json
        └── preferences.json               # Learned preferences
    └── user_456/
        └── ...

# Conversation schema
{
    "session_id": "20250116_001",
    "user_id": "user_123",
    "timestamp": "2025-01-16T10:30:00Z",
    "messages": [...],
    "topics": ["action movies", "2020s cinema"],
    "preferences_learned": {
        "favorite_genres": ["action", "sci-fi"],
        "preferred_platforms": ["netflix"]
    }
}
```

Features to Add:

- 💾 Persist conversations to disk (JSON or SQLite)
- 👤 User-specific history and preferences

- 🔎 Semantic search over past conversations
- 📊 Analytics on user interests
- 💬 Personalized recommendations based on history

Technical Implementation:


- Replace MemorySaver with custom SQLiteCheckpointer
- Add user authentication (see #3)
- Implement conversation summarization for long histories
- Privacy controls (GDPR compliance)

3. 🔒 User Management & API Key Interface

Current Limitation:

- Single shared API keys in `.env`
- No multi-user support
- API costs not attributable to users

Proposed UI:



Implementation:

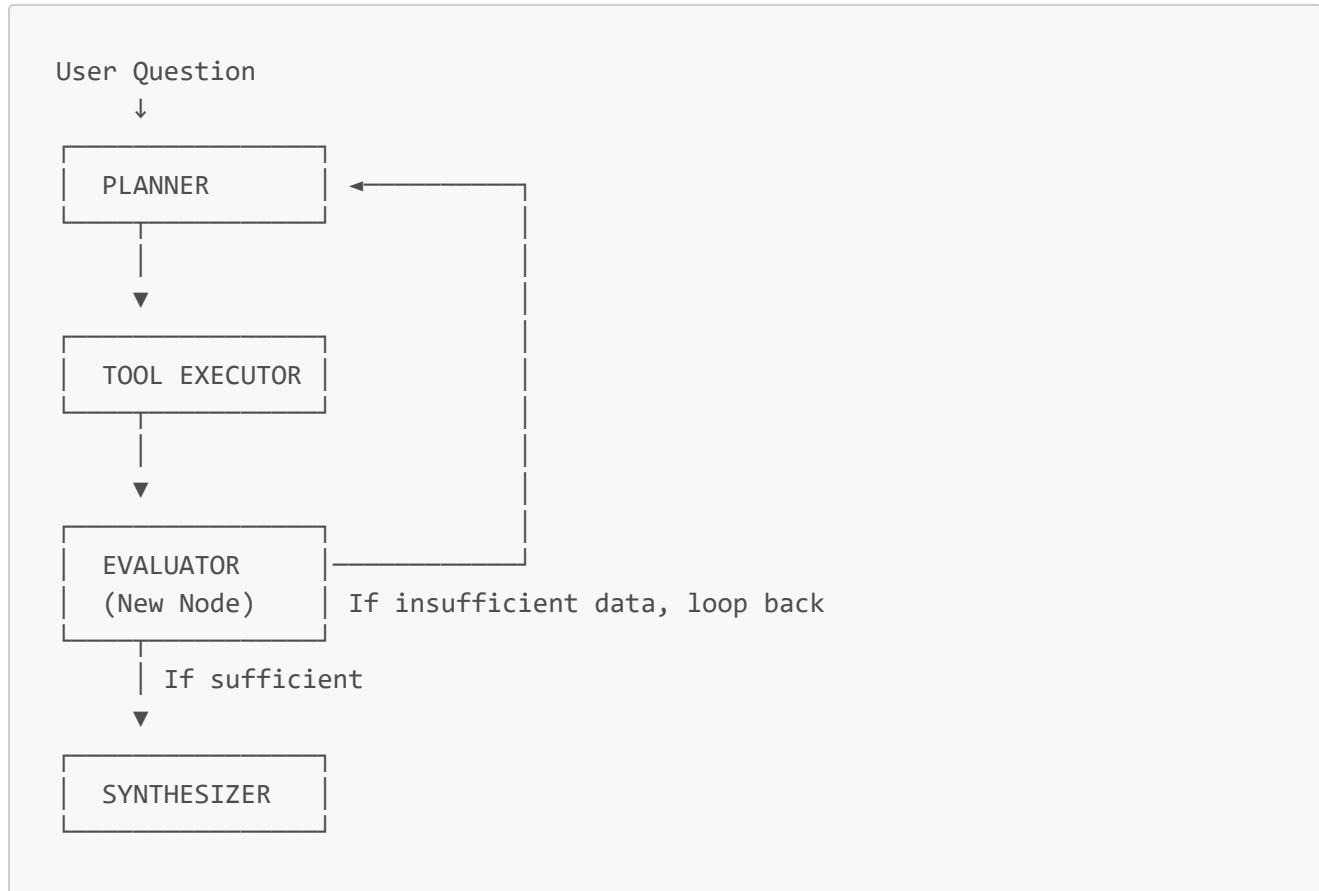
- Streamlit sidebar with settings panel
- Encrypted key storage per user (keyring library)
- Session-based authentication
- Token tracking and cost estimation
- Rate limiting per user

4. 🗃 Workflow Enhancement: Planner Loop

Current Issue:

- Linear workflow: Planner → Tools → Synthesizer → End
- No feedback loop if initial plan was insufficient
- Cannot self-correct or ask for more tools

Proposed Architecture:



Evaluator Node Logic:

```

def evaluator_node(state: AgentState) -> dict:
    """
    Checks if gathered data is sufficient to answer question
    Returns: continue (to synthesizer) or replan (back to planner)
    """
    results_quality = assess_data_completeness(state)

    if results_quality["sufficient"]:
        return {"next": "synthesize"}
    else:
        # Request additional tools
        return {
            "next": "planner",
            "feedback": "Need more data: " + results_quality["missing"]
        }
  
```

Benefits:

- Self-correcting behavior
- More complete answers
- Adaptive tool selection
- Avoids premature responses

5. Embedding Quality Improvements

Current Limitations:

- Similarity scores often low (<50%)
- Movie descriptions are single sentences only
- No chunking strategy
- Basic embedding model (text-embedding-3-small)

Improvement Strategies:

A. Enhance Movie Description Quality (Priority #1)

Currently, we only embed the plot description field from databases (single sentence).

Solution:

Use APIs to enrich our database with way longer movie descriptions.

Expected Impact:

- More contextual embeddings
- Better similarity scores (10-20% improvement)
- Genre/cast matching in semantic search

6. Structured Output Enforcement

Problem:

- JSON parsing errors possible in synthesizer
- Inconsistent response formats
- Difficult to extract structured data

Solution - Pydantic Models:

```
from pydantic import BaseModel, Field
from typing import List, Optional

class MovieResult(BaseModel):
    title: str
    year: int
    rating: Optional[str]
    genres: List[str]
    description: str
    source: str # "sql", "semantic", "omdb"

class QueryResponse(BaseModel):
```

```

    answer: str = Field(..., description="Natural language answer")
    movies: List[MovieResult] = Field(default=[], description="Structured movie
data")
    sources: List[str] = Field(..., description="Sources used")
    confidence: float = Field(..., ge=0, le=1, description="Confidence score")
    follow_up_suggestions: List[str] = Field(default=[], description="Suggested
follow-up questions")

# Use in synthesizer
structured_llm = llm.with_structured_output(QueryResponse)
response = structured_llm.invoke(synthesis_prompt)

```

Benefits:

- Guaranteed valid JSON
- Type safety
- Easier frontend integration
- Better error messages

7. ⚙ Token Optimization

Current Issues:

- Prompts are verbose (500-800 tokens each)
- Full database catalog sent to planner (1000+ tokens)
- Conversation history grows unbounded

Optimization Strategies:

A. Prompt Compression

B. Catalog Summarization

C. Conversation Summarization

D. Lazy Loading

Expected Savings:

- 60-70% reduction in prompt tokens
- ~\$0.02 → \$0.005 per query
- Faster LLM responses

8. 🌐 UI/UX Enhancements

Proposed Features:

- **Results Table View** - Toggle between chat and table display
- **Movie Cards** - Rich display with posters, ratings, cast
- **Query Statistics** - Show token usage, cost, response time

-  **Dark Mode** - Theme switching
-  **Export Results** - Download conversations as JSON/CSV
-  **Voice Input** - Speech-to-text for queries
-  **Multi-language Support** - Full i18n for French/English
-  **Mobile Responsiveness** - Optimize for mobile devices
-  **Keyboard Shortcuts** - Power user features

9. Testing & Quality Assurance

Current Gap: No automated tests

Coverage Goals:

- Unit tests for each node
- Integration tests for workflow
- Performance benchmarks
- Regression tests for common queries

10. Security & Privacy

Enhancements Needed:

-  API key encryption at rest
-  Input sanitization (SQL injection prevention)
-  PII detection and redaction in conversations
-  Audit logging for all queries
-  Rate limiting and abuse prevention
-  HTTPS enforcement in production
-  Content filtering for inappropriate queries

11. Performance & Scalability

Optimization Opportunities:

-  **Async Tool Execution** - Run SQL, Semantic, OMDB in parallel
-  **Result Caching** - Cache common queries (Redis)
-  **Vector Index Optimization** - Use HNSW parameters tuning
-  **Database Indexing** - Add indexes on common query columns
-  **Connection Pooling** - Reuse DB connections
-  **Deployment** - Docker + cloud hosting (AWS/GCP)
-  **CDN Integration** - Cache static assets

12. Analytics & Monitoring

Tracking Metrics:

-  Query latency by tool type
-  Cost per query (token usage)
-  Tool selection accuracy (planner effectiveness)

-  User engagement metrics
-  Error rates and types
-  Most common queries and topics
-  Semantic search quality metrics

Contributors

This project was developed as part of our Master's degree at **Albert School X Mines Paris - PSL**.

Team:

- Vincent Lamy & Alexandre Waerniers

Institution:

- Albert School (Paris, France)
 - Mines Paris - PSL
-

License

This project is licensed under the MIT License - see the LICENSE file for details.

★ If you found this project useful, please consider giving it a star! ★